

```
In [1]: %load_ext autoreload
        %autoreload 2
```

Task 1: Problem definition

Project Description:

The aim of our project is to create a model that allow us to predict the price of hdb using relevant parameters, such as the flat type, size, distance form mrt, etc. We achieved this goal by following the steps below: 1. A suitable dataset with the relevant data and attributes is selected 2. Additional datasets are evaluated and used to enrich the existing dataset. 3. Exploratory data analysis and visualisation are done where data is cleaned and prepared 4. Multiple models are trained and evaluated using RMSE (Root-mean-square error) and MAE (Mean absolute error). 5. The selected model is analysed where the feature importance is evaluated

Potential use cases:

1. The trained model can be deployed on real estate website such as PropertyGuru, 99.co, where buyers and sellers are able to get the most optimal price estimation before listing or buying.
2. This model can also be used by HDB to real estate agent to find a suitable price when consulting their clients on calculating the valuation of the client's property.

Task 2: Data collection/curation

We have collected the data we required for our HDB sales price prediction task primary from data.gov.sg.

HDB resale price :

https://beta.data.gov.sg/collections/189/datasets/d_8b84c4ee58e3cfc0ece0d773c8ca6abc/view

Besides, we also enriched our dataset using additional data from kaggle as follows:

HDB coordinate : <https://www.kaggle.com/datasets/mylee2009/singapore-postal-code-mapper> **MRT coordinate** : <https://www.kaggle.com/datasets/shengjunlim/singapore-mrt-lrt-stations-with-coordinates>

The files have been downloaded locally and renamed accordingly.

Task 3 & 4: Data preparation & Exploratory data analysis and visualization

Step 1: Import the packages and load resale data

```
In [2]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns

%matplotlib inline

data = pd.read_csv("ResaleflatpricesbasedonregistrationdatefromJan2017onwards.csv")
hdb_locations = pd.read_csv("sg_zipcode_mapper_utf.csv", delimiter=',')
mrt_locations = pd.read_csv("mrt.csv")
```

Step 2: Data cleaning and EDA

```
In [3]: data.info()

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 177572 entries, 0 to 177571
Data columns (total 11 columns):
 #   Column                Non-Null Count  Dtype  
---  -
 0   month                 177572 non-null object  
 1   town                  177572 non-null object  
 2   flat_type             177572 non-null object  
 3   block                 177572 non-null object  
 4   street_name           177572 non-null object  
 5   storey_range          177572 non-null object  
 6   floor_area_sqm        177572 non-null float64  
 7   flat_model            177572 non-null object  
 8   lease_commence_date   177572 non-null int64  
 9   remaining_lease       177572 non-null object  
 10  resale_price           177572 non-null float64  
dtypes: float64(2), int64(1), object(8)
memory usage: 14.9+ MB
```

```
In [4]: data.head()
```

Out[4]:

	month	town	flat_type	block	street_name	storey_range	floor_area_sqm	flat_model
0	2017-01	ANG MO KIO	2 ROOM	406	ANG MO KIO AVE 10	10 TO 12	44.0	Improved
1	2017-01	ANG MO KIO	3 ROOM	108	ANG MO KIO AVE 4	01 TO 03	67.0	New Generation
2	2017-01	ANG MO KIO	3 ROOM	602	ANG MO KIO AVE 5	01 TO 03	67.0	New Generation
3	2017-01	ANG MO KIO	3 ROOM	465	ANG MO KIO AVE 10	04 TO 06	68.0	New Generation
4	2017-01	ANG MO KIO	3 ROOM	601	ANG MO KIO AVE 5	01 TO 03	67.0	New Generation

The data is already cleaned, it does not contain null value, therefore data cleaning is not required.

Extract year from the data and filter data to only include data from 2019 to 2023, because we only want to predict the price of the flat based on the past 4 years of data.

```
In [5]: data['month'] = pd.to_datetime(data['month'])
data['year'] = list(map(lambda x: x.year, data['month']))

data = data[data['year'] >= 2019]
data = data[data['year'] != 2024]

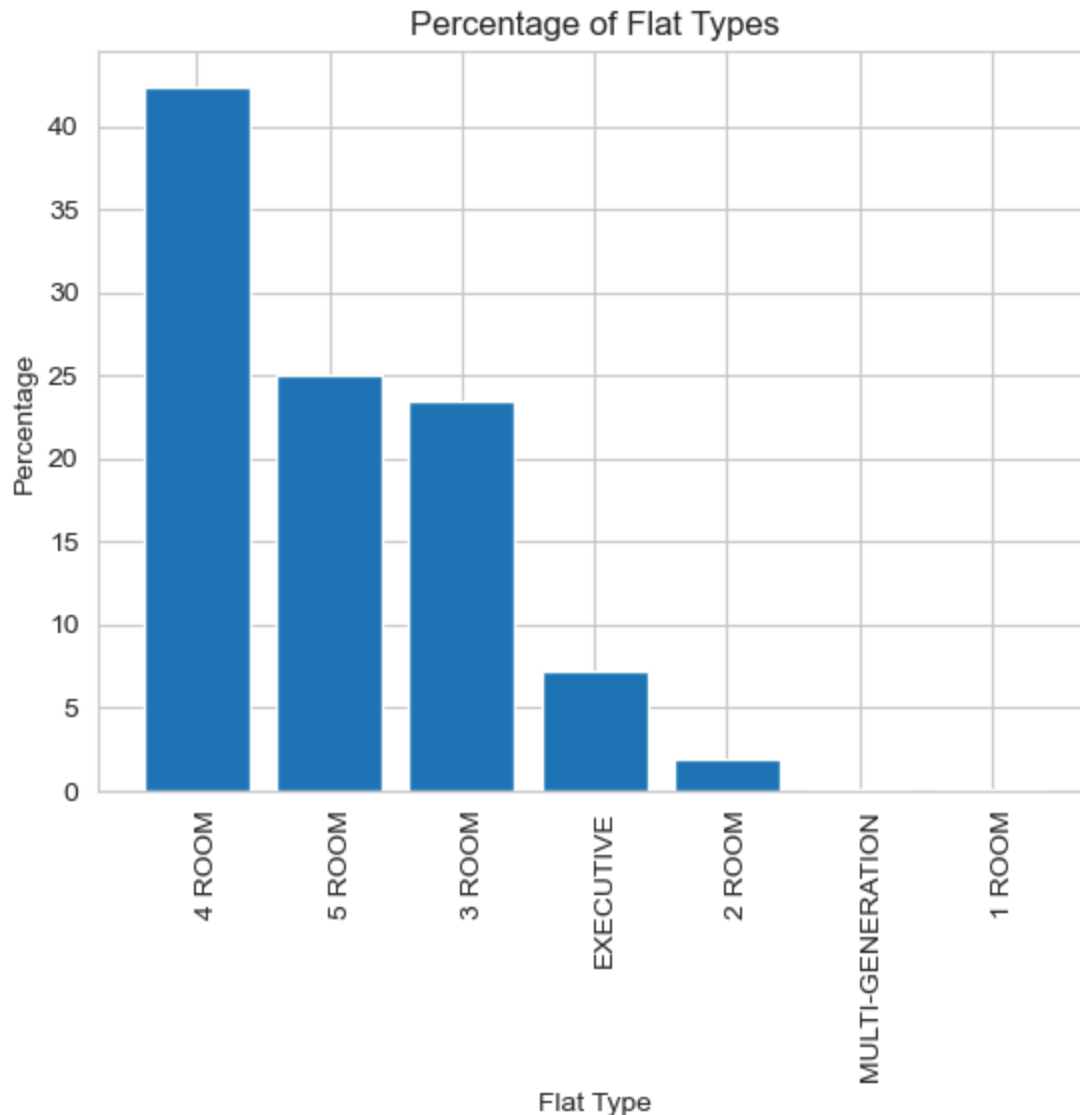
data['price_per_sqm'] = data['resale_price'] / data['floor_area_sqm']
data['remaining_lease'] = (99 + data['lease_commence_date']) - data['year']
```

```
In [6]: flat_type_counts = data['flat_type'].value_counts(normalize=True) * 100 # This calculates the percentage of each flat type

# Create a bar plot using Matplotlib
plt.bar(flat_type_counts.index, flat_type_counts.values)

plt.xlabel('Flat Type')
plt.ylabel('Percentage')
plt.title('Percentage of Flat Types')
plt.xticks(rotation=90, ha='left', rotation_mode='default')
plt.show()

# # Show the plot
# plt.show()
```



From the count plot above we can see that there are very little one room, two room and multi-generatin flat(below 5%). therefore we are going to drop these data for our model training.

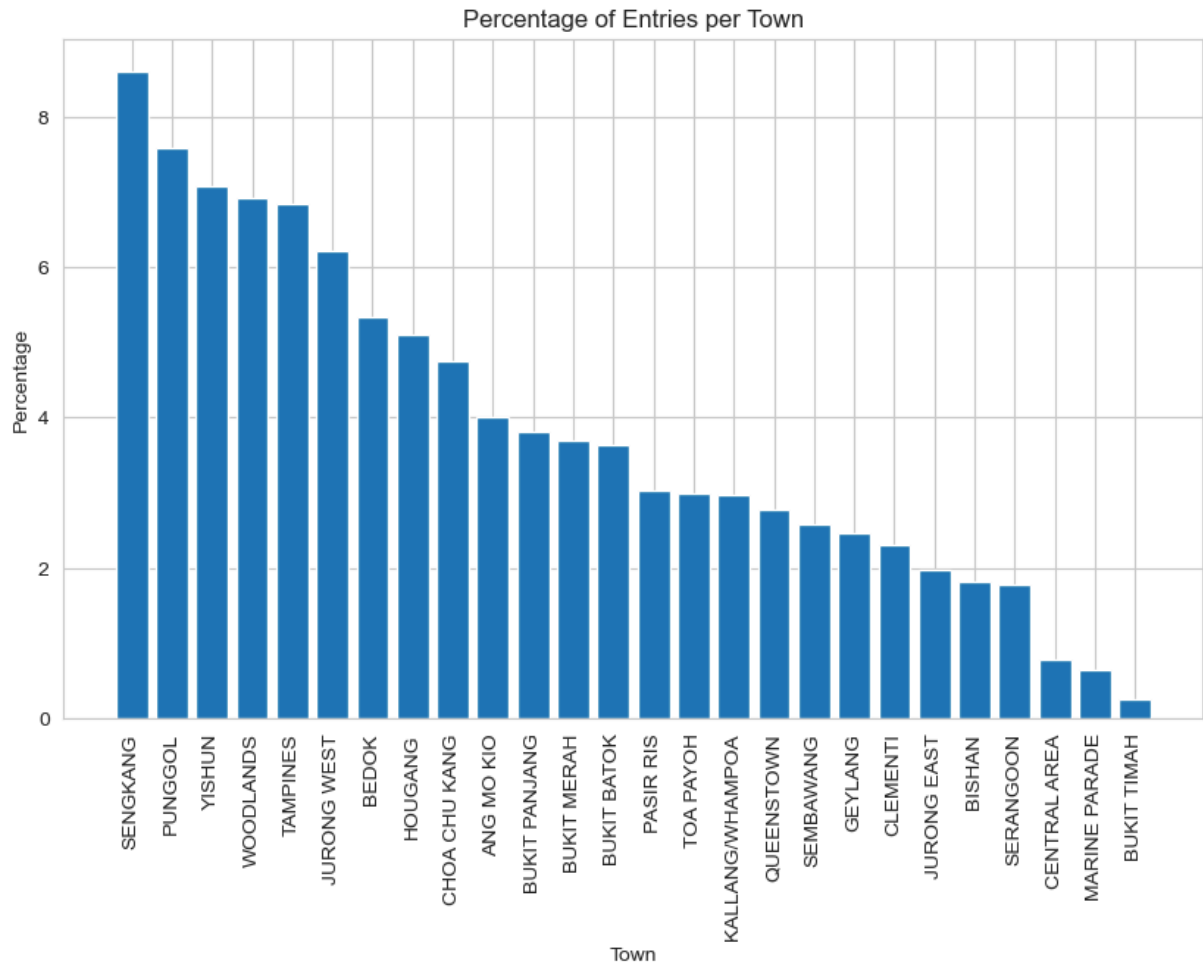
```
In [7]: data = data[data['flat_type'] != "1 ROOM"]
data = data[data['flat_type'] != "2 ROOM"]
data = data[data['flat_type'] != "MULTI-GENERATION"]
```

```
In [8]: town_counts = data['town'].value_counts() # Get the counts for each town
total_counts = town_counts.sum() # Total number of entries
town_percent = (town_counts / total_counts) * 100 # Convert counts to percentage

# Plotting with Matplotlib
plt.figure(figsize=(10, 6)) # Optional: Adjust the size of the figure
plt.bar(town_percent.index, town_percent.values) # Create a bar chart

plt.xlabel('Town') # Label for the x-axis
plt.ylabel('Percentage') # Label for the y-axis
plt.title('Percentage of Entries per Town') # Title of the plot
```

```
plt.xticks(rotation=90) # Rotate the x-axis labels for better visibility if needed
plt.show()
```



From the count plot above we can see that there are very little ales recoed in Central Area, Bukit Timah and Marine Parade(below 1%).therefore we are going to drop these data for our model training.

```
In [9]: data = data[data['town'] != "CENTRAL AREA"]
data = data[data['town'] != "BUKIT TIMAH"]
data = data[data['town'] != "MARINE PARADE"]
```

Load the HDB location file, and join with the resale dataframe. This will allow us to calculate the distance from the nearest mrt for each HDB blocks.


```
In [10]: hdb_locations.rename(columns={"latitude": "lat", "longtitude": "lng",
                                     "blk_no": "block", "RD_name": "street_name"},
                              inplace=True)
hdb_locations = hdb_locations[hdb_locations['postal.1'] != 189677]
hdb_locations = hdb_locations[hdb_locations['postal.1'] != 469277]
hdb_locations = hdb_locations[hdb_locations['postal'] != 90027]
hdb_locations = hdb_locations[["postal", "lat", "lng", "block", "street_name", "pos
imputed_data = pd.merge(data, hdb_locations, how='left',
```

```
on=['block', 'street_name'])
imputed_data.dropna(inplace=True)
```

```
In [11]: imputed_data.head()
```

```
Out[11]:
```

	month	town	flat_type	block	street_name	storey_range	floor_area_sqm	flat_model
0	2019-01-01	ANG MO KIO	3 ROOM	225	ANG MO KIO AVE 1	01 TO 03	67.0	New Generation
1	2019-01-01	ANG MO KIO	3 ROOM	174	ANG MO KIO AVE 4	01 TO 03	60.0	Improved
2	2019-01-01	ANG MO KIO	3 ROOM	440	ANG MO KIO AVE 10	04 TO 06	67.0	New Generation
3	2019-01-01	ANG MO KIO	3 ROOM	174	ANG MO KIO AVE 4	10 TO 12	61.0	Improved
4	2019-01-01	ANG MO KIO	3 ROOM	637	ANG MO KIO AVE 6	01 TO 03	68.0	New Generation



Data Enrichment

We will try to enrich the data by calculating the distance of the HDB from MRT stations as we think distance to mrt station might be a good predictor

Here we will clean the mrt_locations data to remove the unnecessary data

```
In [12]: mrt_locations.head()
```

Out[12]:

	Unnamed: 0	OBJECTID	STN_NAME	STN_NO	geometry	Latitude	Longitude
0	0	1	EUNOS MRT STATION	EW7	POINT (103.9032524667383 1.319778951553637)	1.319779	103.903252
1	1	2	CHINESE GARDEN MRT STATION	EW25	POINT (103.7325967380734 1.342352820874744)	1.342353	103.732597
2	2	3	KHATIB MRT STATION	NS14	POINT (103.8329799077383 1.417383370153547)	1.417383	103.832980
3	3	4	KRANJI MRT STATION	NS7	POINT (103.7621654109002 1.425177698770448)	1.425178	103.762165
4	4	5	REDHILL MRT STATION	EW18	POINT (103.816816670149 1.289562726402453)	1.289563	103.816817

```
In [13]: mrt_locations.drop(columns=['Unnamed: 0', 'OBJECTID',
                                     'geometry', "STN_NO", ], inplace=True)
```

```
In [14]: mrt_locations.rename(columns={"Latitude": "lat", "Longitude": "lng"}, inplace=True)
```

```
In [15]: mrt_locations.head()
```

Out[15]:

	STN_NAME	lat	lng
0	EUNOS MRT STATION	1.319779	103.903252
1	CHINESE GARDEN MRT STATION	1.342353	103.732597
2	KHATIB MRT STATION	1.417383	103.832980
3	KRANJI MRT STATION	1.425178	103.762165
4	REDHILL MRT STATION	1.289563	103.816817

This is the formula to convert from euclidean distance to earth distance

Formula taken from <https://byteshiva.medium.com/navigate-the-world-a-guide-to-calculating-distances-between-points-on-earth-using-python-and-the-9f3c5c856203>

```
In [16]: def calculate_distance(x, y):
          R = 6371 # Radius of the Earth in kilometers

          lat_a, lon_a = x[0], x[1]
          lat_b, lon_b = y[0], y[1]
```

```

# Convert Latitude and Longitude to radians
lat_a, lon_a, lat_b, lon_b = np.radians([lat_a, lon_a, lat_b, lon_b])

# Calculate the difference in Latitude and Longitude
dlat = lat_b - lat_a
dlon = lon_b - lon_a

# Apply the Haversine formula
a = np.sin(dlat / 2) ** 2 + np.cos(lat_a) * np.cos(lat_b) * np.sin(dlon / 2) ** 2
c = 2 * np.arctan2(np.sqrt(a), np.sqrt(1 - a))
distance = R * c

return distance

```

Here we import the KNN package to calculate the distance between each HDB block and the nearest mrt

```

In [17]: from sklearn.neighbors import KNeighborsClassifier

nbrs = KNeighborsClassifier(n_neighbors=1, algorithm='ball_tree',
                           metric=calculate_distance).fit(mrt_locations[['lat', 'lon'],
                               mrt_locations['STN_NAME']

imputed_data['mrt_dist'] = nbrs.kneighbors(imputed_data[['lat', 'lng']])[0]

imputed_data['nearest_mrt'] = nbrs.predict(imputed_data[['lat', 'lng']])

```

```

In [18]: imputed_data.head()

```

```

Out[18]:

```

	month	town	flat_type	block	street_name	storey_range	floor_area_sqm	flat_model
0	2019-01-01	ANG MO KIO	3 ROOM	225	ANG MO KIO AVE 1	01 TO 03	67.0	New Generation
1	2019-01-01	ANG MO KIO	3 ROOM	174	ANG MO KIO AVE 4	01 TO 03	60.0	Improved
2	2019-01-01	ANG MO KIO	3 ROOM	440	ANG MO KIO AVE 10	04 TO 06	67.0	New Generation
3	2019-01-01	ANG MO KIO	3 ROOM	174	ANG MO KIO AVE 4	10 TO 12	61.0	Improved
4	2019-01-01	ANG MO KIO	3 ROOM	637	ANG MO KIO AVE 6	01 TO 03	68.0	New Generation

```

In [19]: UniqueNames = imputed_data['town'].unique()

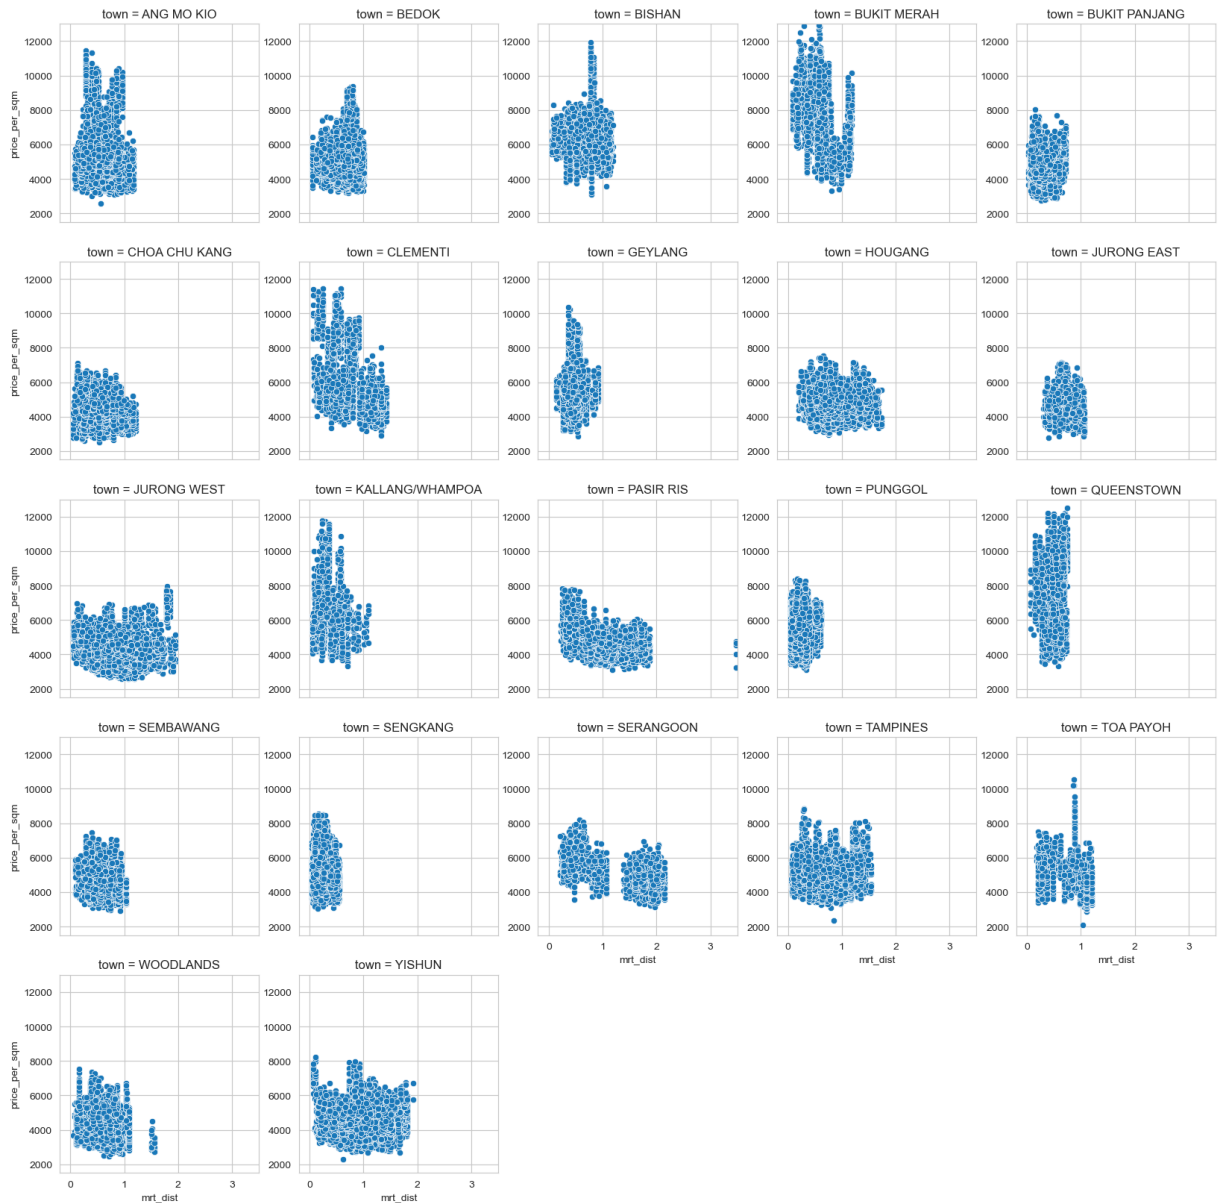
```



```
#create a data frame dictionary to store your data frames
DataFrameDict = {elem: pd.DataFrame() for elem in UniqueNames}

for key in DataFrameDict.keys():
    DataFrameDict[key] = imputed_data[imputed_data.town == key]
f = plt.figure(figsize=(20, 20))
for i in range(len(UniqueNames)):
    column = i // 5
    row = i % 5
    axis = f.add_subplot(5, 5, i + 1)
    dataset = UniqueNames[i]
    title = "town = " + dataset
    axis.scatter(DataFrameDict[dataset]['mrt_dist'],
                 DataFrameDict[dataset]['price_per_sqm'],
                 color='#1978b9',
                 edgecolors='white',
                 linewidth=0.5)
    axis.set_xlim([-0.2, 3.5])
    axis.set_ylim([1500, 13000])
    axis.set_xticks([0, 1, 2, 3])
    axis.set_title(title)

    if i >= 17:
        axis.set(xlabel='mrt_dist')
    else:
        axis.tick_params(colors='black', which='major', axis='x', labelcolor='white')
    if row == 0:
        axis.set(ylabel='price_per_sqm')
```



From the chart of price_per_sqm versus mrt_dist plot we generated for each town, we observed there is not a strong correlation between price and mrt distance for most of the towns. Thus, the mrt_dist might be a useful predictor when training our model.

Visualisation

To visualise the hot spot where of hdb location and their distance to the mrt stations, we have plot the graph below resembling the Singapore map

```
In [20]: map_town_data = pd.DataFrame({"Latitude": imputed_data["lat"],
                                       "Longitude": imputed_data["lng"],
                                       "Town": imputed_data["town"]})

train_map_data = pd.DataFrame({"Latitude": mrt_locations["lat"],
                               "Longitude": mrt_locations["lng"]})

custom_palette = ['#1f77b4', '#ff7f0e', '#2ca02c', '#d62728', '#9467bd', '#8c564b',
```

```


        '#17becf', '#9edae5', '#393b79', '#5254a3', '#6b6ecf', '#9c9ede',
        '#cedb9c', '#8c6d31', '#bd9e39', '#e7ba52']

sns.set(style="whitegrid")
plt.figure(figsize=(10, 6))

sns.scatterplot(x="Longitude", y="Latitude", hue="Town", data=map_town_data, palette=
sns.scatterplot(x="Longitude", y="Latitude", data=train_map_data, facecolor="red",
plt.legend(title="Towns in Singapore", loc="center left", bbox_to_anchor=(1, 0.5))

plt.xlabel("Longitude")
plt.ylabel("Latitude")
plt.title("Map of HDB resale flats with MRT & LRT stations from the period of 2019
plt.show()


```



```

In [21]: corr_matrix = imputed_data.select_dtypes(include=['float64', 'int64']).corr()
sns.heatmap(corr_matrix)
plt.show()

```



From the heatmap shown above, we theorised that floor_area_sqm and remaining_lease will be a good predictor of resale price, while mrt_dist will not contribute much in predicting the resale_price

Task 5: Modeling and visualization

Feature selection for model training

To reduce the complexity of the model and prevent model overfitting, we have discarded:

- 1) features which are too sparse(block & street name)
- 2) features which are similar(lease commence date) - the value is already calculated and represented by remaining lease

```
In [22]: data = imputed_data[
    ["month", "town", "flat_type", "storey_range", "floor_area_sqm", "flat_model",
     "resale_price"]]
```

```
In [23]: data.to_csv('output.csv', index=False)
```

Step 3: Modelling

1) Linear Regression Model

We train a linear regression model as a baseline where we will improve on it using more complicated models which utilises techniques such as cross validation, bagging, stacking and ensembling.

```
In [24]: data = pd.read_csv('output.csv')
```

```
In [25]: # Data preparation
from sklearn.model_selection import train_test_split
from sklearn.metrics import mean_squared_error, r2_score, mean_absolute_error

# Split data into predictors and target
X = data.drop('resale_price', axis=1)
X = X.drop('month', axis=1)
y = data['resale_price']

# Split the dataset into training (90%) and test (10%) sets
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.1, random_sta
```

Preprocessing for categorical features

Before piping the data into the regression model, we perform 1 hot encoding to transform Categorical data (which are text data) into one hot encoding to match the input format of the Linear Regressor

```
In [26]: from sklearn.linear_model import LinearRegression
from sklearn.preprocessing import OneHotEncoder
from sklearn.compose import ColumnTransformer
from sklearn.pipeline import Pipeline

# Define columns that need encoding
categorical_features = ['town', 'flat_type', 'storey_range', 'flat_model']

# OneHotEncoder for categorical data within a ColumnTransformer
preprocessor = ColumnTransformer(
    transformers=[
        ('cat', OneHotEncoder(handle_unknown='ignore'), categorical_features)
    ],
    remainder='passthrough' # Include all other columns in the model unchanged
)

# Create a pipeline that includes preprocessing and the linear model
LR_model = Pipeline(steps=[
    ('preprocessor', preprocessor),
    ('regressor', LinearRegression())
])

# Train the model
LR_model.fit(X_train, y_train)

# Predict on the test data
y_pred = LR_model.predict(X_test)
```

```
# Evaluate the model
mse = mean_squared_error(y_test, y_pred)
r2 = r2_score(y_test, y_pred)
mae = mean_absolute_error(y_test, y_pred)
rmse = (mean_squared_error(y_test, y_pred))**0.5

print('RMSE: {:.2f}'.format(rmse))
print('MAE: {:.2f}'.format(mae))
print('r2: {:.2f}'.format(r2))
```

RMSE: 79845.39

MAE: 64428.02

r2: 0.77

2) Decision Tree

With the baseline identified, we proceeded to experiment with more complex models such as DecisionTreeRegressor. The DecisionTreeRegressor is a versatile machine learning model used for regression tasks that predicts continuous outcomes by splitting data into increasingly homogeneous subsets. It constructs a tree-like model of decisions based on the values of the input features, making it highly interpretable and easy to visualize.

By using the decision tree regressor, we are able to improve the RMSE from 79845 to 54105, with a 47% improvement.

```
In [27]: from sklearn.tree import DecisionTreeRegressor

DT_model = Pipeline(steps=[
    ('preprocessor', preprocessor),
    ('regressor', DecisionTreeRegressor(random_state=42))
])

# Train the model
DT_model.fit(X_train, y_train)

# Predict on the test data
y_pred = DT_model.predict(X_test)

# Evaluate the model

r2 = r2_score(y_test, y_pred)
mae = mean_absolute_error(y_test, y_pred)
rmse = (mean_squared_error(y_test, y_pred))**0.5

print('RMSE: {:.2f}'.format(rmse))
print('MAE: {:.2f}'.format(mae))
print('r2: {:.2f}'.format(r2))
```

RMSE: 54104.85

MAE: 36034.01

r2: 0.90

3) Hyperparameter tuning on Decision Tree Regressor

Then, the model is finetuned by optimizing the parameters that govern the training process. In the cell below, we test various combinations of hyperparameters to determine the most effective settings for the Decision Tree Regressor including, max_depth, min_samples_split and min_samples_leaf.

It shows significant improvement (6.1%) with the RMSE decreasing from 54105 to 50984.

```
In [28]: from sklearn.model_selection import GridSearchCV

# Define the parameter grid
param_grid = {
    'regressor__max_depth': [None, 10, 20, 30], # Maximum number of levels in each
    'regressor__min_samples_split': [2, 10, 20], # Minimum number of samples requi
    'regressor__min_samples_leaf': [1, 5, 10] # Minimum number of samples required
}

grid_search = GridSearchCV(estimator=DT_model, param_grid=param_grid, scoring='neg_

# Fit GridSearchCV
grid_search.fit(X_train, y_train)

# Get the best parameters and best score
print("Best parameters:", grid_search.best_params_)
print("Best score:", grid_search.best_score_)

best_model = grid_search.best_estimator_

# Predict on the test data
y_pred = best_model.predict(X_test)

# Evaluate the model

r2 = r2_score(y_test, y_pred)
mae = mean_absolute_error(y_test, y_pred)
rmse = (mean_squared_error(y_test, y_pred))**0.5

print('RMSE: {:.2f}'.format(rmse))
print('MAE: {:.2f}'.format(mae))
print('r2: {:.2f}'.format(r2))
```

Fitting 5 folds for each of 36 candidates, totalling 180 fits

Best parameters: {'regressor__max_depth': None, 'regressor__min_samples_leaf': 1, 'regressor__min_samples_split': 10}

Best score: -2702001284.5219145

RMSE: 50984.10

MAE: 34885.60

r2: 0.91

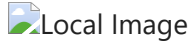
Note: The below part is experimental and requires the installation of additional packages!

```
In [37]: !pip install autogluon
# !pip install -U scikit-learn==1.4.0
```

4) AutoML

By referring to the paper [AutoGluon-Tabular](#), we have tried the best performing AutoML model which is able to automatically select relevant models based on our modelling task (regression) such as KNN, RandomForest, LightGBM, XGBoost. Then, a final model is trained using ensembling where the output of the most performant models are weighted.

By doing so, we are able to reduce the RMSE from 50984 to 25716.



```
In [30]: # Initialisation
from autogluon.tabular import TabularDataset, TabularPredictor

# Train test split
train, test = train_test_split(data, test_size=0.1, random_state=0)
train_data = TabularDataset(train)
test_data = TabularDataset(test)
```

```
In [31]: label = 'resale_price'
save_path = './AutogluonModels/small'
try:
    predictor = TabularPredictor.load(save_path)
except:
    # Retraining will take around 60 seconds
    presets = ['good_quality', 'optimize_for_deployment']
    predictor = TabularPredictor(label=label, path=save_path).fit(train_data, prese
```

```
In [32]: predictor.evaluate(test_data)
```

```
Out[32]: {'root_mean_squared_error': -28650.780324820356,
'mean_squared_error': -820867213.2211133,
'mean_absolute_error': -20450.017244982548,
'r2': 0.9707724441181854,
'pearsonr': 0.9852810733672099,
'median_absolute_error': -14724.5625}
```

```
In [33]: results = predictor.evaluate(test_data)

print('RMSE: {:.2f}'.format(-results['root_mean_squared_error']))
print('MAE: {:.2f}'.format(-results['mean_absolute_error']))
print('r2: {:.2f}'.format(results['r2']))
```

```
RMSE: 28650.78
MAE: 20450.02
r2: 0.97
```

Feature importance

By calculating the feature importance scores for the given model via permutation importance, we found that the town and flat_type of the hdb flat have the highest

importance when predicting the resale price.

In [34]: `predictor.feature_importance(test_data)`

Out[34]:

	importance	stddev	p_value	n	p99_high	p99_low
town	106068.756376	1667.738286	7.331617e-09	5	109502.652558	102634.860194
flat_type	85959.846495	1352.574806	7.353609e-09	5	88744.816715	83174.876276
month	59488.467628	473.545541	4.817929e-10	5	60463.504443	58513.430813
remaining_lease	57255.777302	1188.005105	2.222941e-08	5	59701.896234	54809.658370
floor_area_sqm	43517.972184	1176.790124	6.410315e-08	5	45940.999314	41094.945054
flat_model	24563.499687	502.472489	2.100038e-08	5	25598.097487	23528.901888
mrt_dist	18048.738272	713.642695	2.926930e-07	5	19518.138445	16579.338099
storey_range	16548.739849	652.336897	2.891413e-07	5	17891.910536	15205.569161

Demo

With the trained model, we are able to test out our model for our potential users (potential buyer searching for HDB flat near NUS on PropertyGuru):

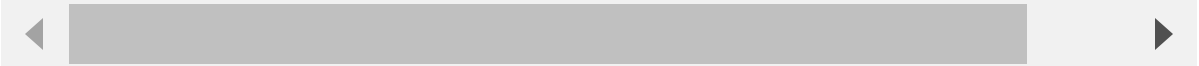
```
In [35]: # Search criteria: a relatively new house near NUS
month = '2024-04-24'
town = 'CLEMENTI'
flat_type = '3 ROOM'
storey_range = '10 TO 12'
floor_area_sqm = 90
flat_model = 'New Generation'
remaining_lease = 90
mrt_distance = 1.0

# hdb_search = data.sample(1)

hdb_search = pd.DataFrame([[month, town, flat_type, storey_range, floor_area_sqm, f
hdb_search
```

Out[35]:

	month	town	flat_type	storey_range	floor_area_sqm	flat_model	remaining_lease
0	2024-04-24	CLEMENTI	3 ROOM	10 TO 12	90	New Generation	90



In [36]: `predictor.predict(hdb_search)`

Out[36]: 0 629813.625
 Name: resale_price, dtype: float32

Conclusion

In the data preparation stage, we selected the most relevant data by filtering out the outliers. We also enriched the data by calculating the distance of mrt to hdb.

In the EDA stage, we visualised the data to have a better understanding of the data and plotted a correlation matrix in a heatmap. We also preprocessed the data by applying one hot encoding on categorical data.

In the modelling stage, we saw that even simple model like linear regression can get a relatively good prediction. However, we found out that more complex model just as Decision Tree Regression is able to get a better results, and can be further enhance using hyperparameter tuning. Lastly, we also tried out the state of the art model which utilises ensembling for best model performance.

In the feature importance analysis, we also confirmed our hypothesis when calculating the correlation matrix, where mrt_distance does not really contribute in the prediction of sales price.

In conclusion, we have successfully trained a highly accurate sales price prediction model with a 5.1% error (RMSE = 25716, mean = 507,654) after going through the steps of EDA, data enrichment, data preparation and modelling.